

Data Communication & Networks

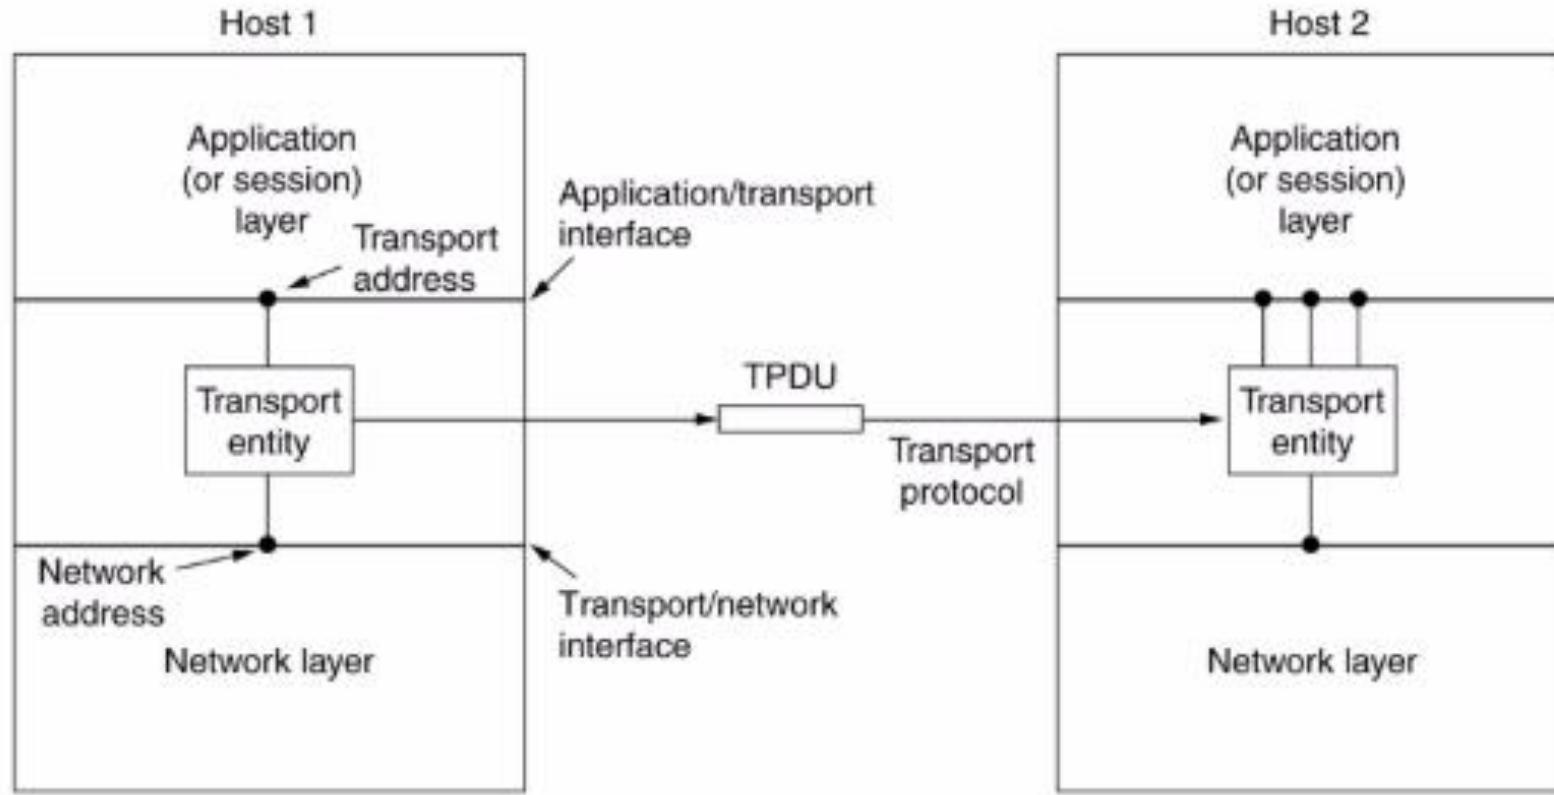
Chapter 6: Transport Layer (Textbook)

Transport Layer

The Transport Service

- Services Provided to the Upper Layers
- Transport Service Primitives
- Berkeley Sockets
- An Example of Socket Programming:
 - An Internet File Server

Services Provided to the Upper Layer



The network, transport, and application layers.

Transport Service Primitives

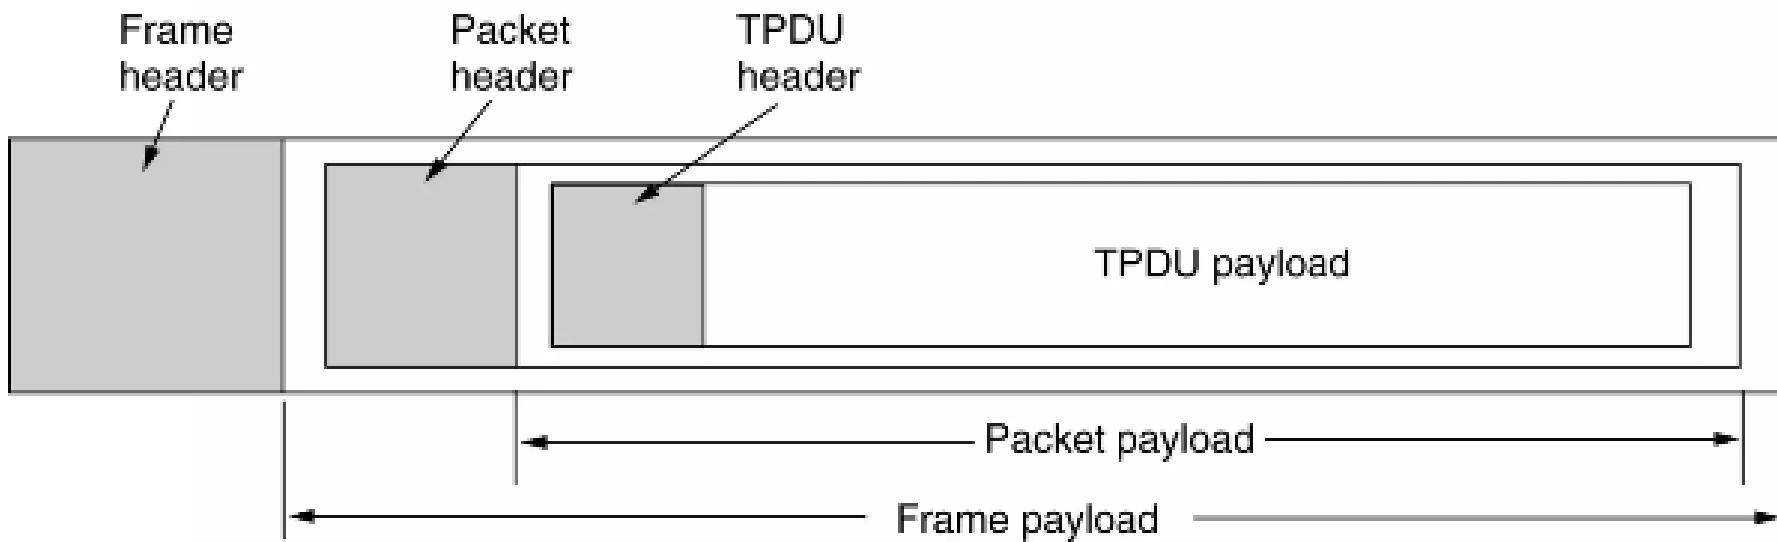
Transport Service Primitives

| Primitive | Packet sent | Meaning |
|------------|--------------------|--------------------------------------------|
| LISTEN | (none) | Block until some process tries to connect |
| CONNECT | CONNECTION REQ. | Actively attempt to establish a connection |
| SEND | DATA | Send information |
| RECEIVE | (none) | Block until a DATA packet arrives |
| DISCONNECT | DISCONNECTION REQ. | This side wants to release the connection |

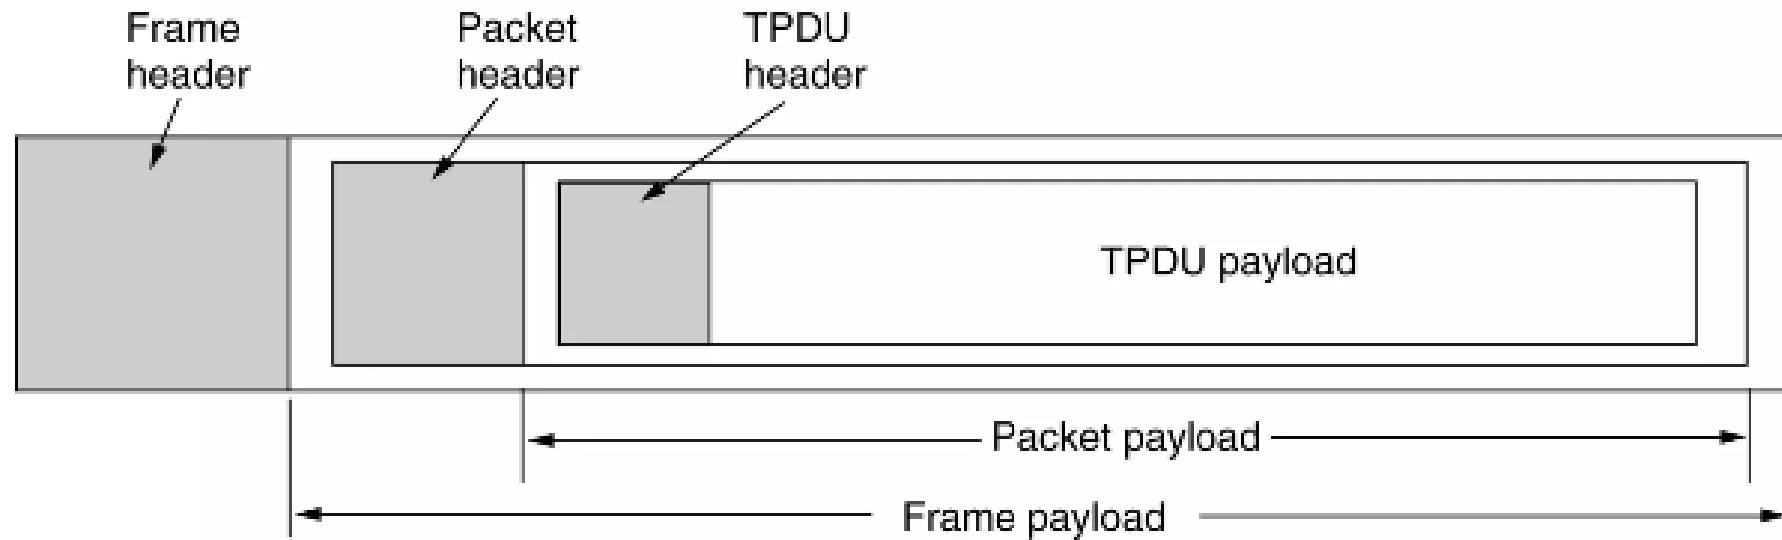
The primitives for a simple transport service.

- A quick note on terminology is now in order. For lack of a better term, we will use the term segment for messages sent from transport entity to transport entity. TCP, UDP and other Internet protocols use this term. Some older protocols used the ungainly name TPDU (Transport Protocol Data Unit)

Transport Service Primitives

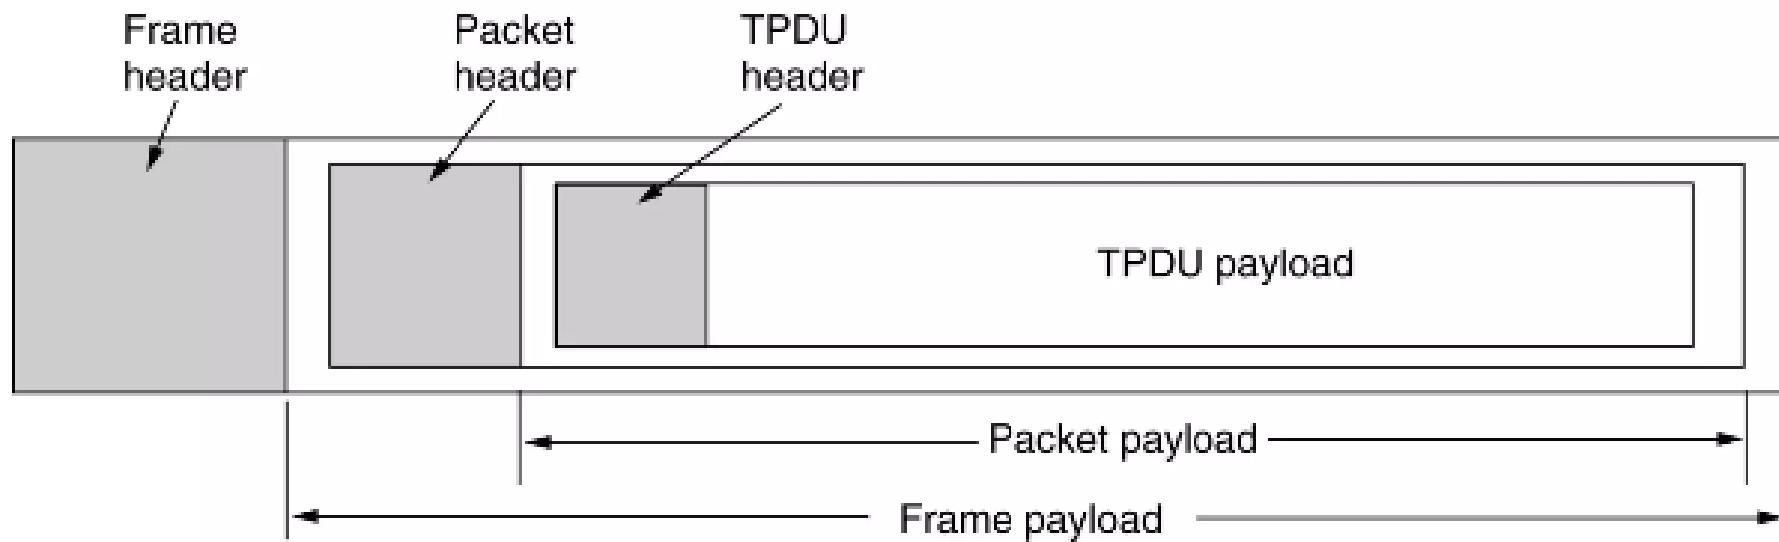


Transport Service Primitives



- Getting back to our client-server example, the client's CONNECT call causes a CONNECTION REQUEST segment to be sent to the server. When it arrives, the transport entity checks to see that the server is blocked on a LISTEN (i.e., is ready to handle requests). If so, it then unblocks the server and sends a CONNECTION ACCEPTED segment back to the client. When this segment arrives, the client is unblocked and the connection is established.

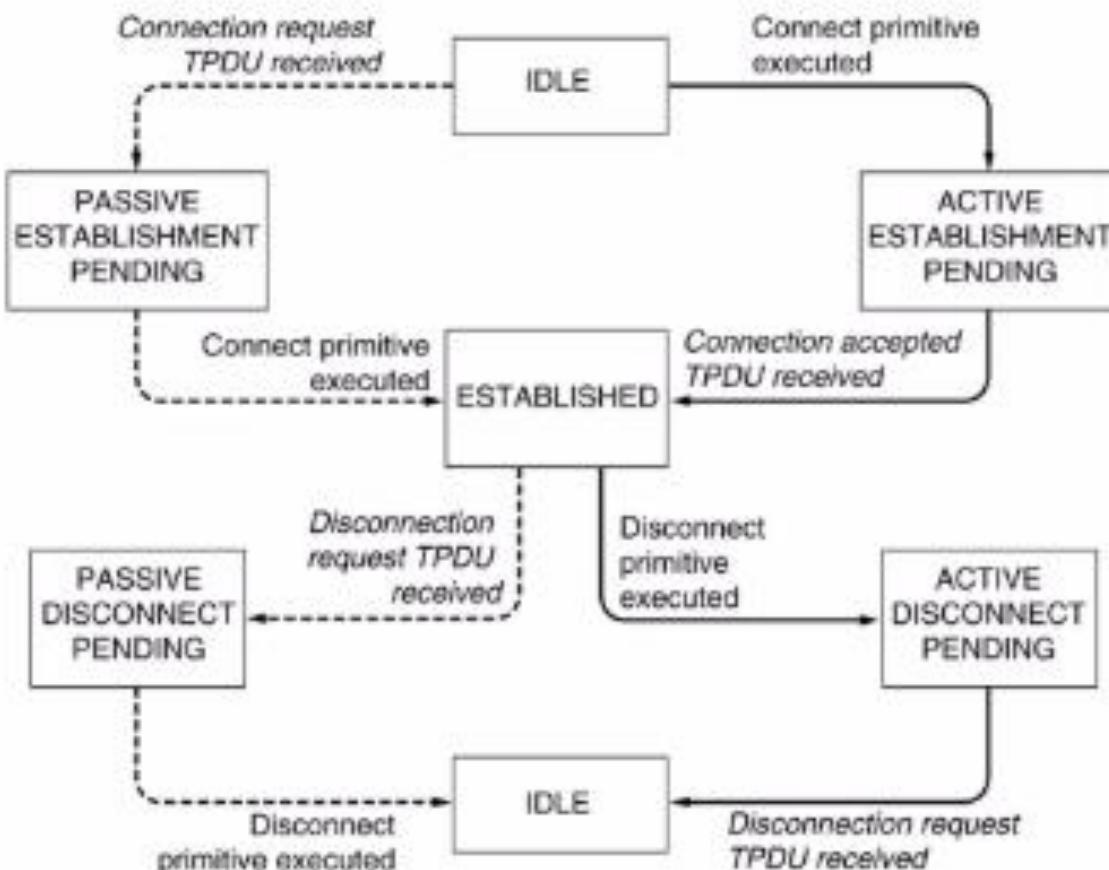
Transport Service Primitives



- Every data segment is acknowledged in Transport Layer.

Transport Service Primitives

Transport Service Primitives (3)



Transport Protocol Elements

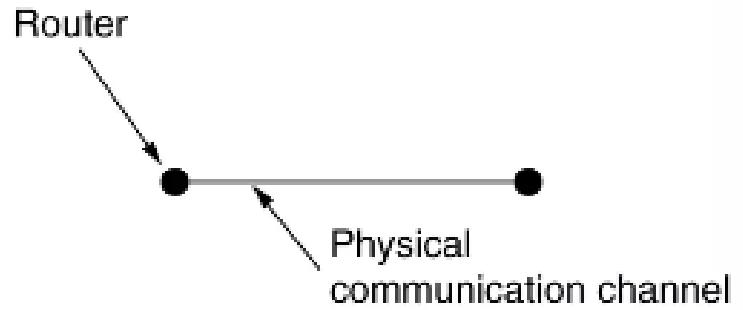
Elements of Transport Protocols

ddressing
action
g *

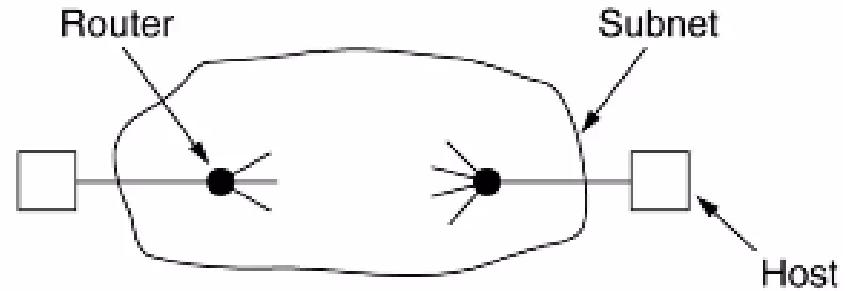
- Addressing
- Connection Establishment
- Connection Release
- Flow Control and Buffering
- Multiplexing
- Crash Recovery

Transport Protocol Elements

Transport Protocol



(a)

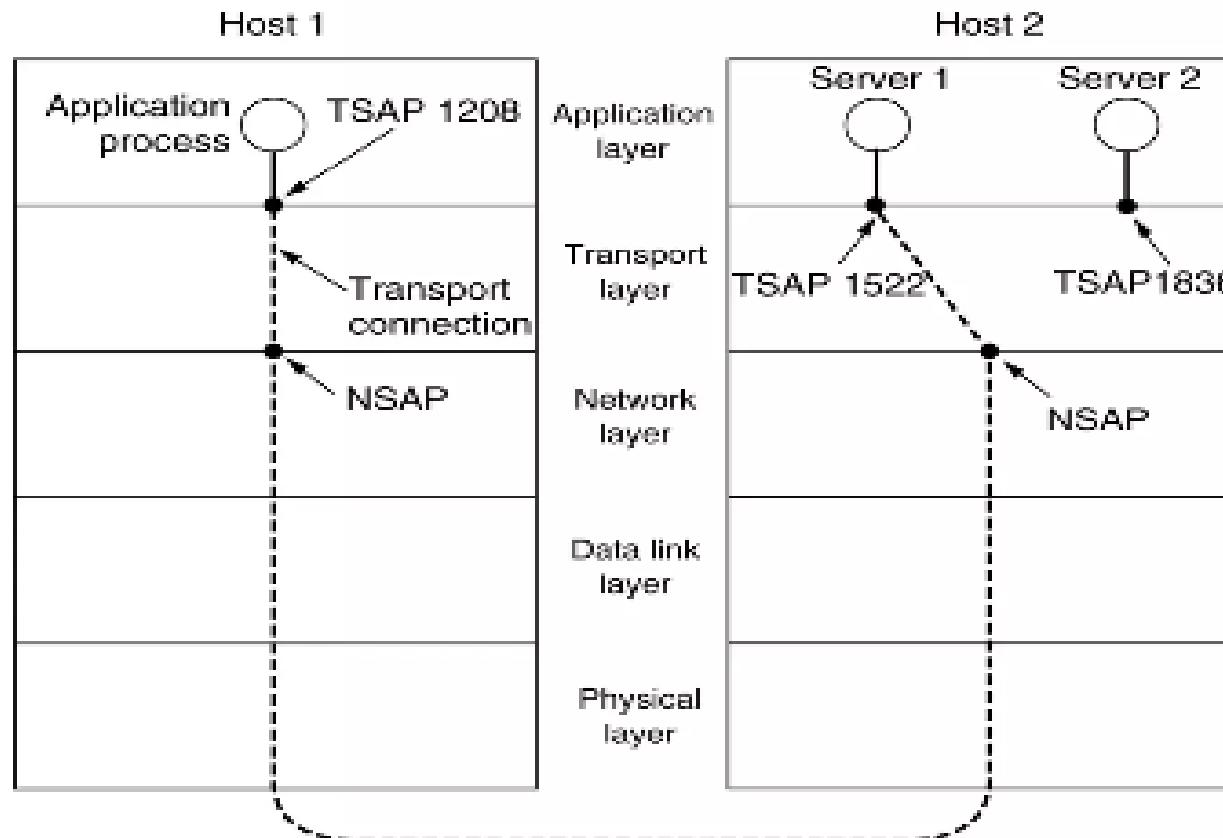


(b)

- (a) Environment of the data link layer.
- (b) Environment of the transport layer.

Transport Protocol Elements

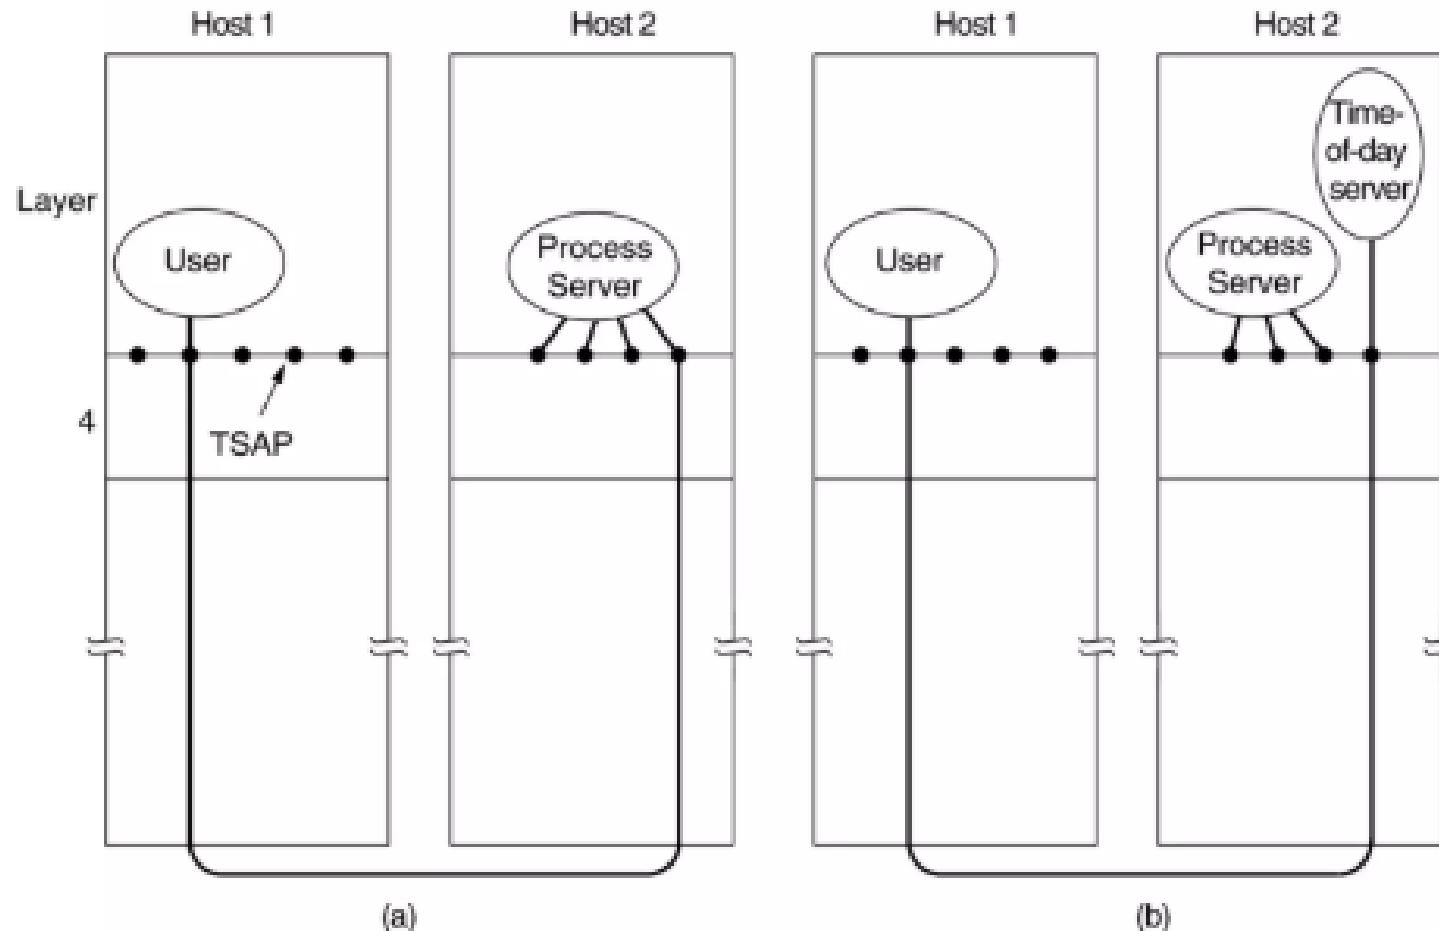
Addressing



TSAPs, NSAPs and transport connections.

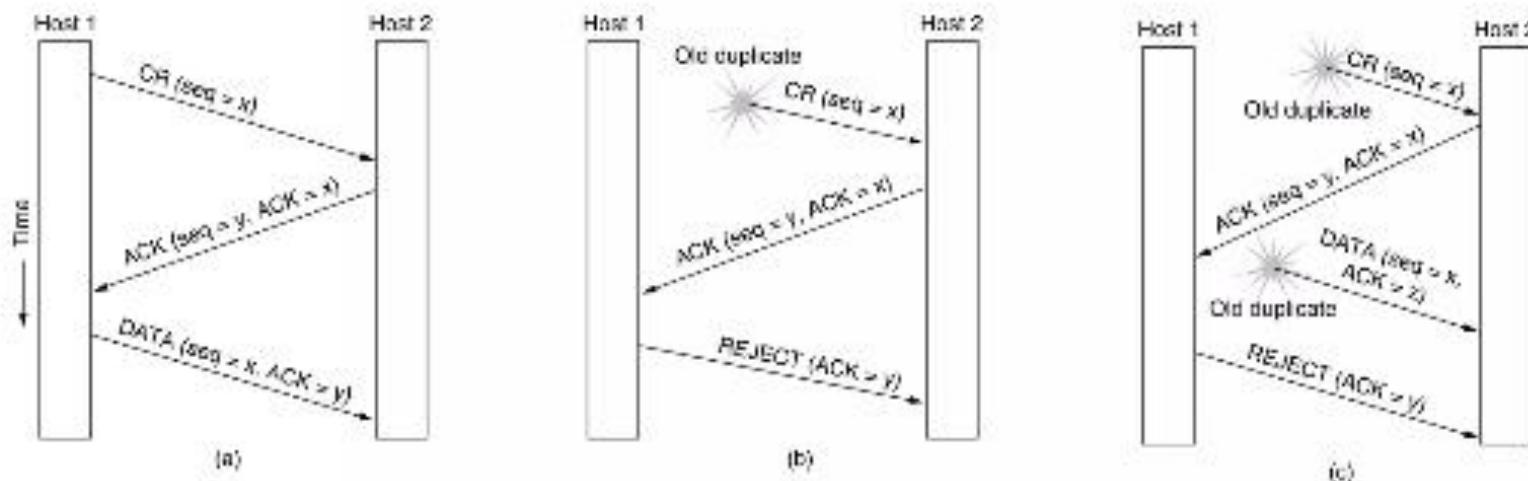
Transport Protocol Elements

Connection Establishment



Transport Protocol Elements

Connection Establishment (3)

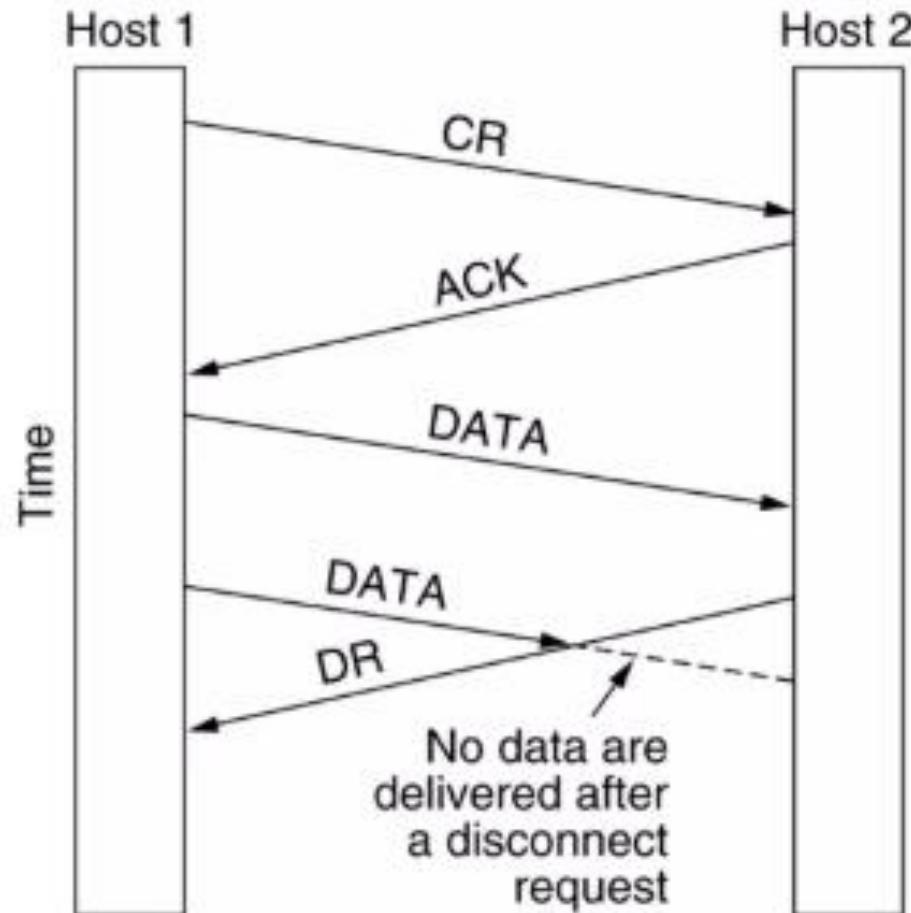


Three protocol scenarios for establishing a connection using a three-way handshake. CR denotes CONNECTION REQUEST

- (a) Normal operation,
- (b) Old CONNECTION REQUEST appearing out of nowhere.
- (c) Duplicate CONNECTION REQUEST and duplicate ACK.

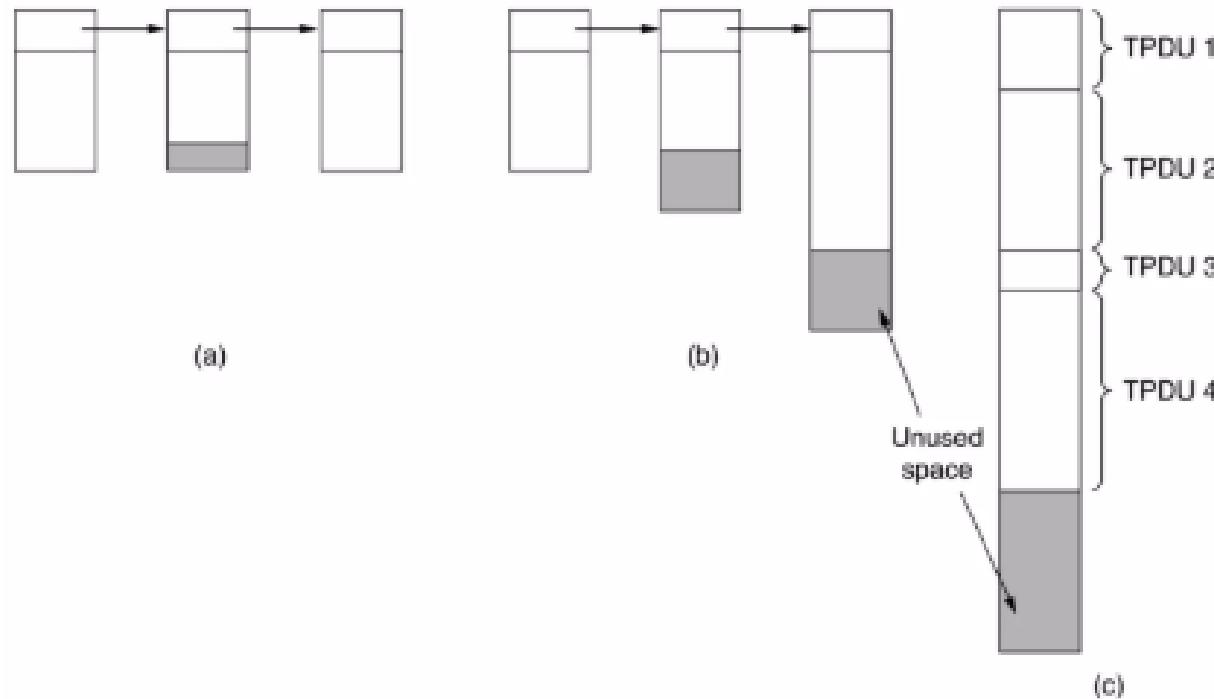
Transport Protocol Elements

Connection Release



Transport Protocol Elements

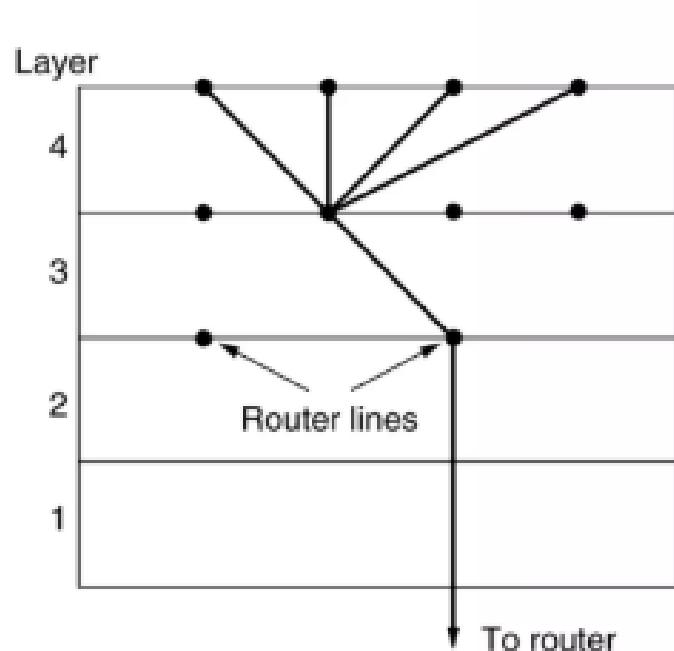
Flow Control and Buffering



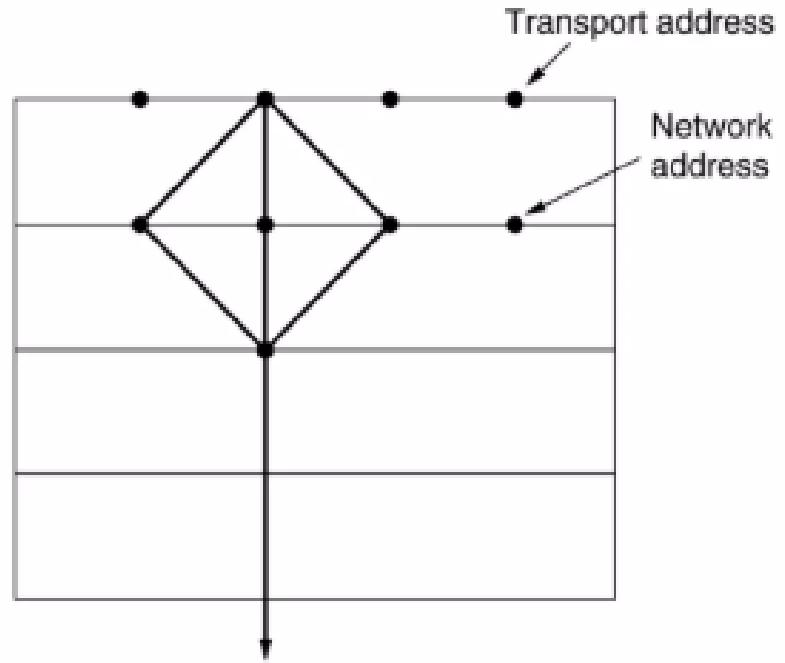
- (a) Chained fixed-size buffers.
- (b) Chained variable-sized buffers.
- (c) One large circular buffer per connection.

Transport Protocol Elements

Multiplexing



(a)



(b)

(a) Upward multiplexing. (b) Downward multiplexing.

Transport Protocol Elements

Crash Recovery

| | | Strategy used by receiving host | | |
|-------------------|--|---------------------------------|------|-----------------------|
| | | First ACK, then write | | First write, then ACK |
| | | AC(W) | AWC | C(AW) |
| Always retransmit | | OK | DUP | OK |
| Never retransmit | | LOST | OK | LOST |
| Retransmit in S0 | | OK | DUP | LOST |
| Retransmit in S1 | | LOST | OK | OK |
| | | C(WA) | W AC | WC(A) |
| | | OK | DUP | DUP |
| | | LOST | OK | OK |
| | | LOST | DUP | OK |
| | | OK | OK | DUP |

OK = Protocol functions correctly

DUP = Protocol generates a duplicate message

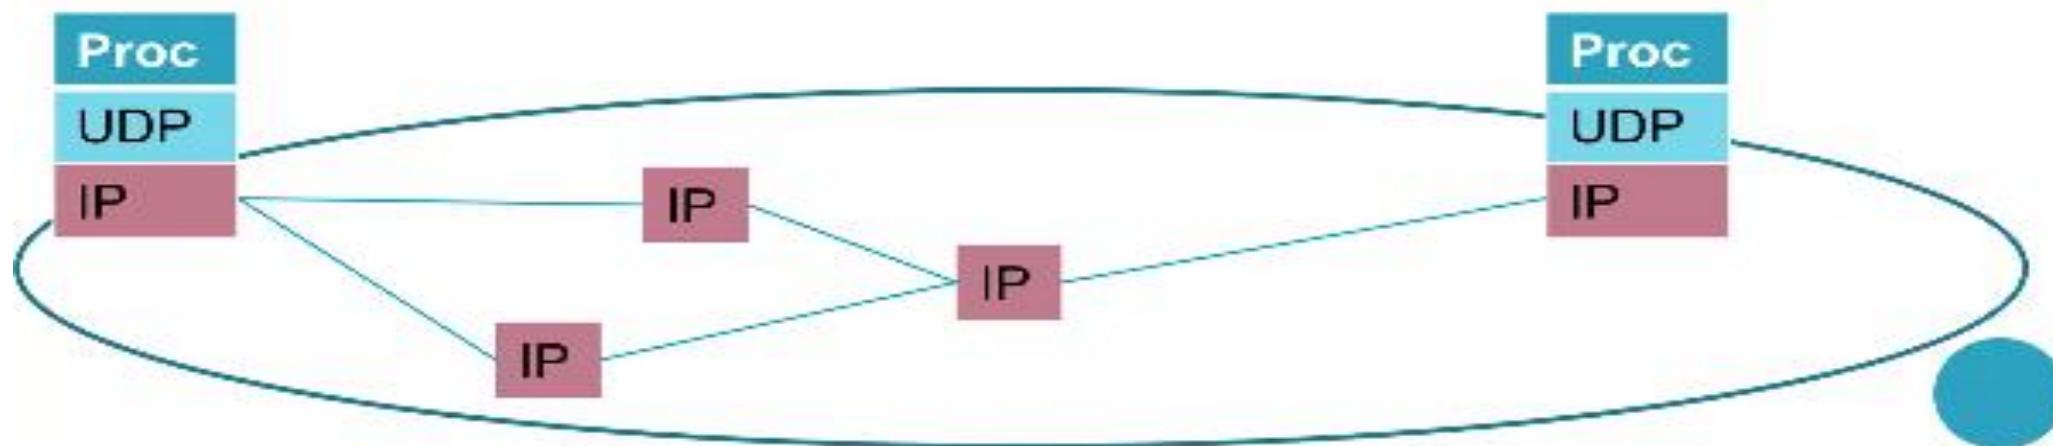
LOST = Protocol loses a message

Different combinations of client and server strategy.

UDP Protocol

Layer 4 (Transport Layer) protocol

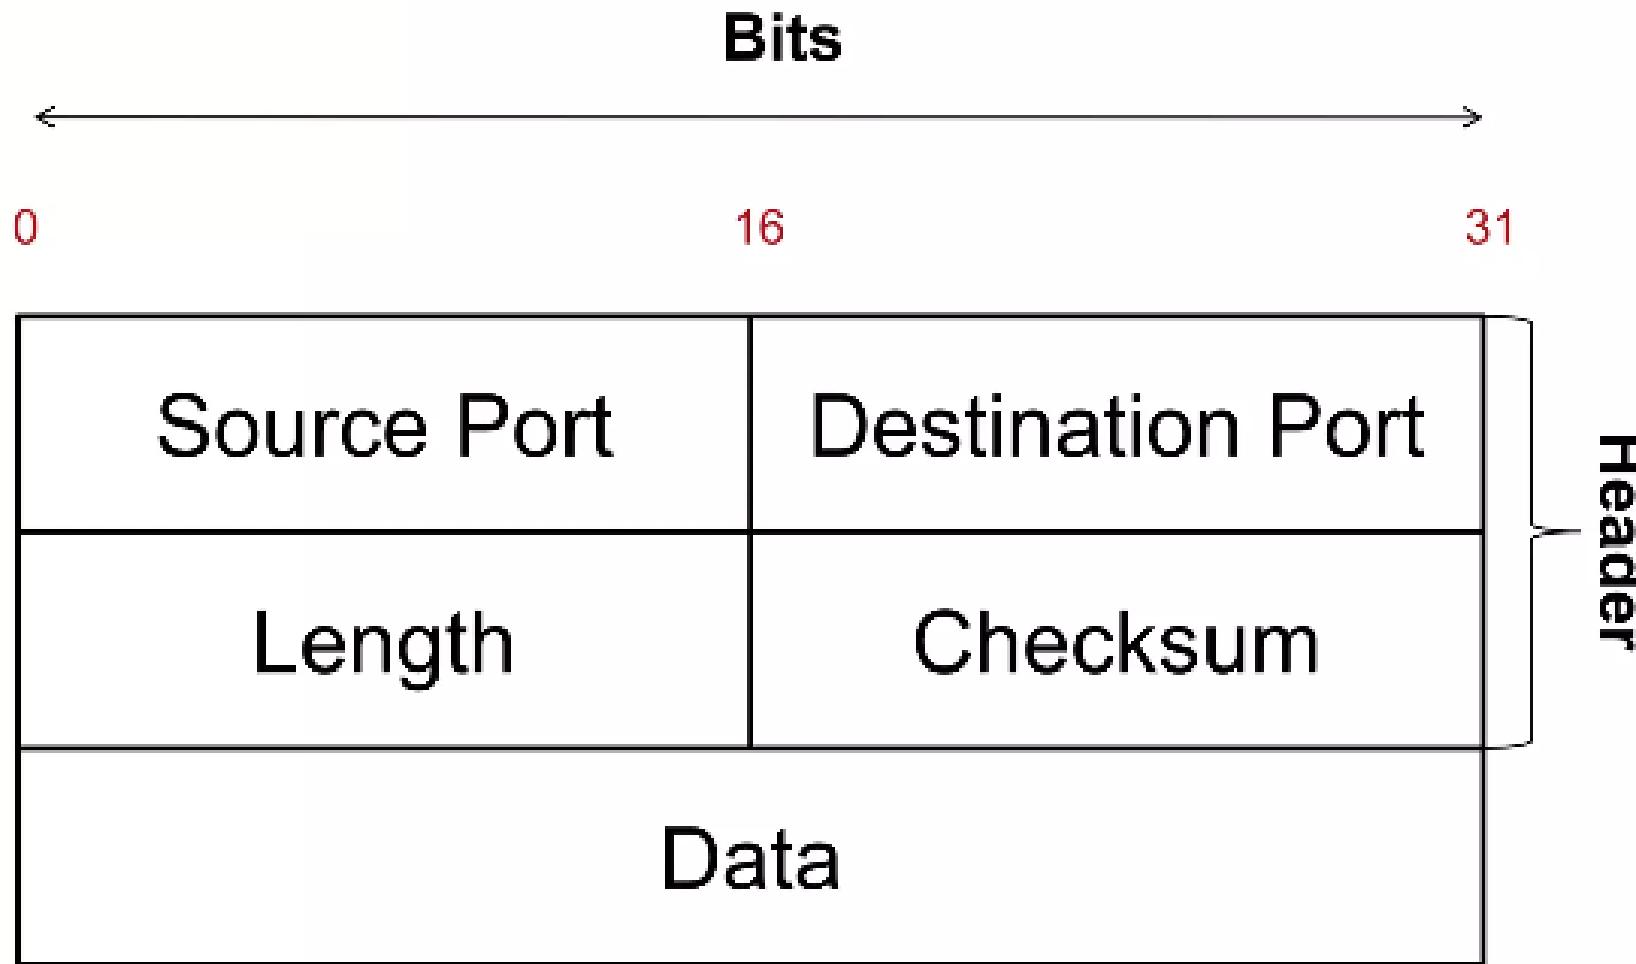
- Runs over IP
- Unreliable i.e. Best Effort service
 - UDP segments may be:
 - lost
 - delivered out of order
- Connectionless
 - no handshaking between UDP sender, receiver
 - each UDP segment handled independently of others



Provides a *lossy* connection (data may vanish).

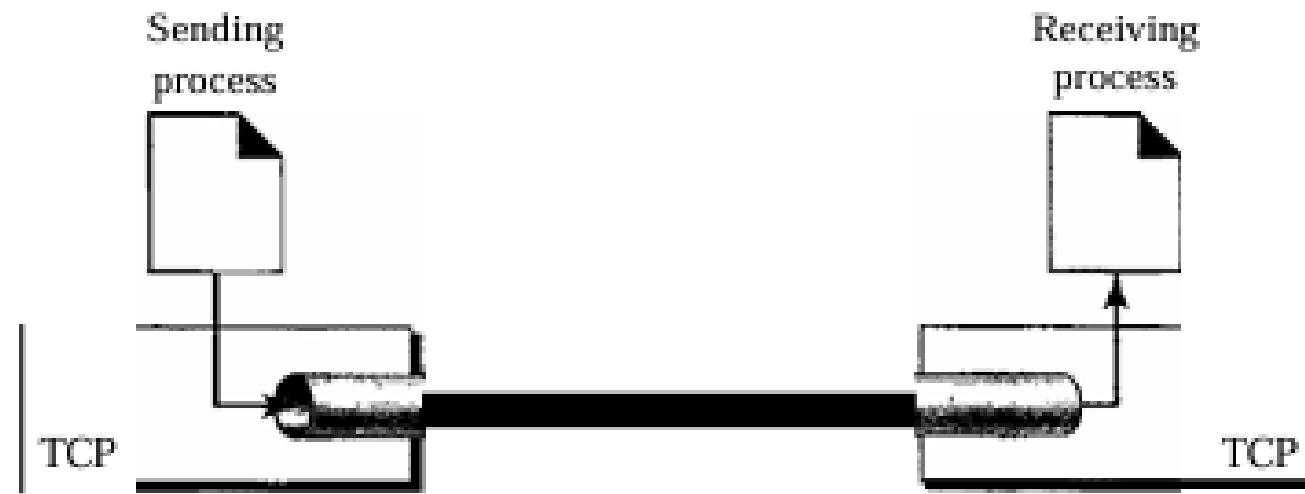
UDP Protocol

8 Bytes Header + Variable Payload



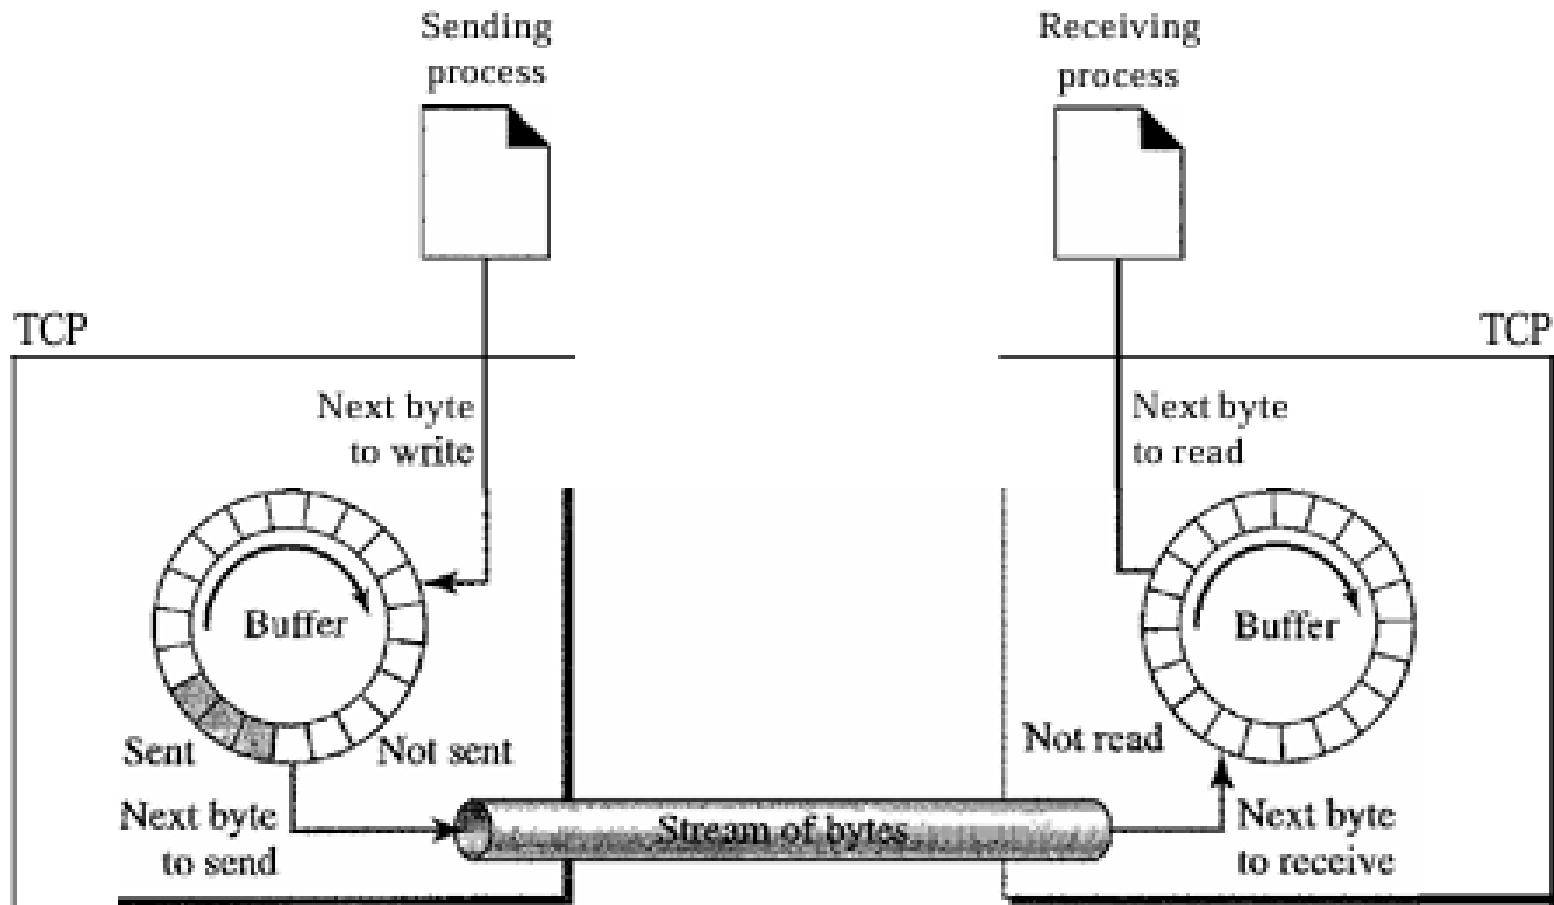
TCP Protocol (Behrouz Ferozan)

Figure 23.13 *Stream delivery*



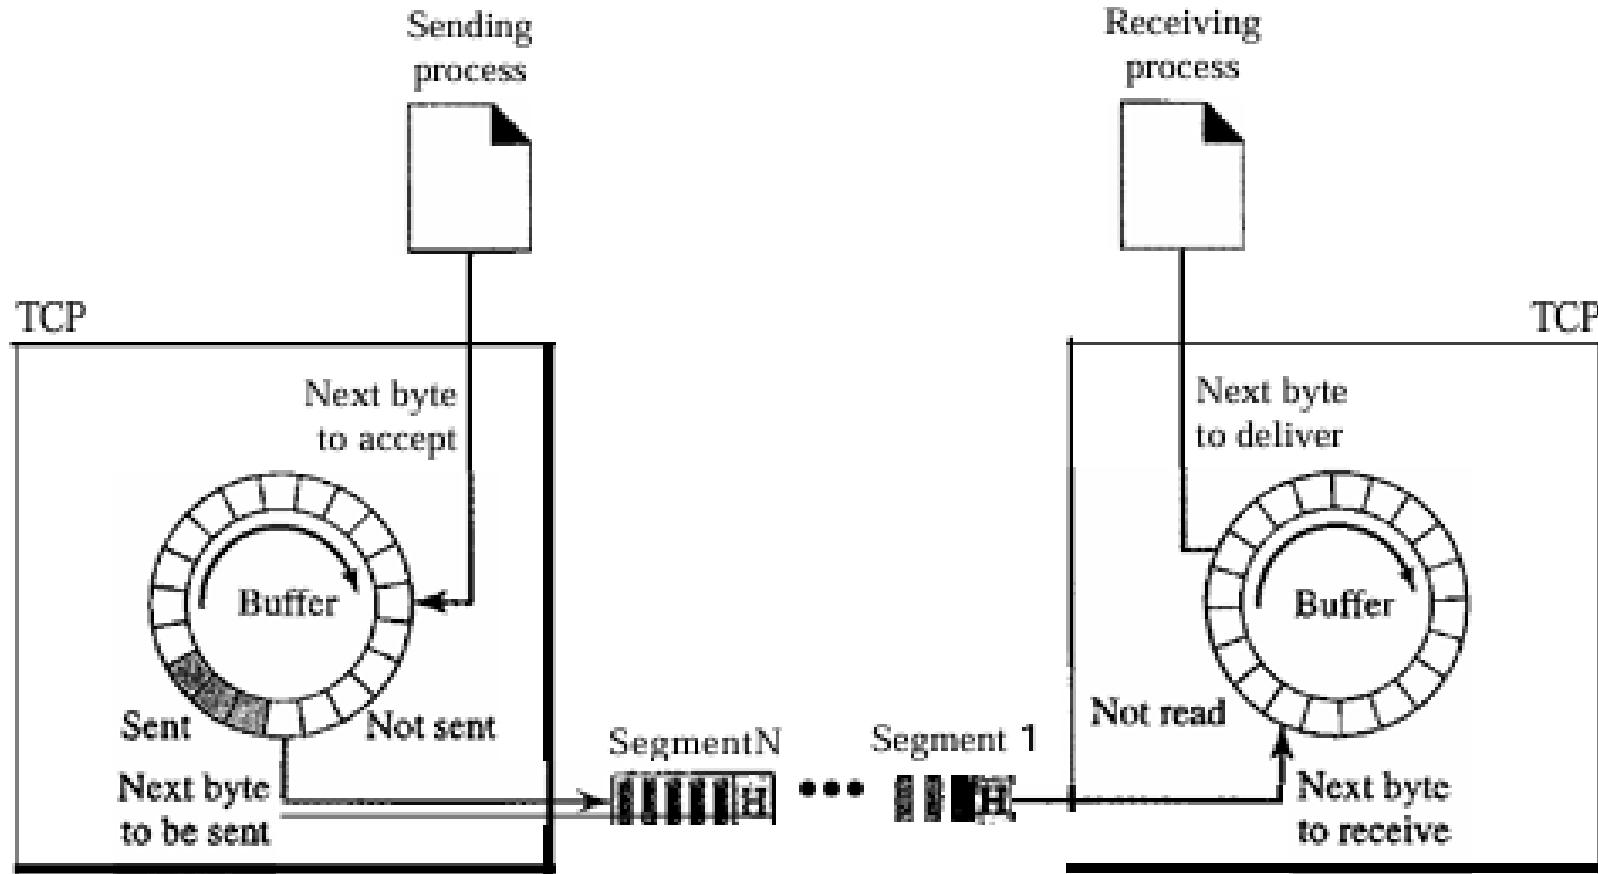
TCP Protocol (Behrouz Ferozan)

Figure 23.14 *Sending and receiving buffers*



TCP Protocol (Behrouz Ferozan)

Figure 23.15 *TCP segments*

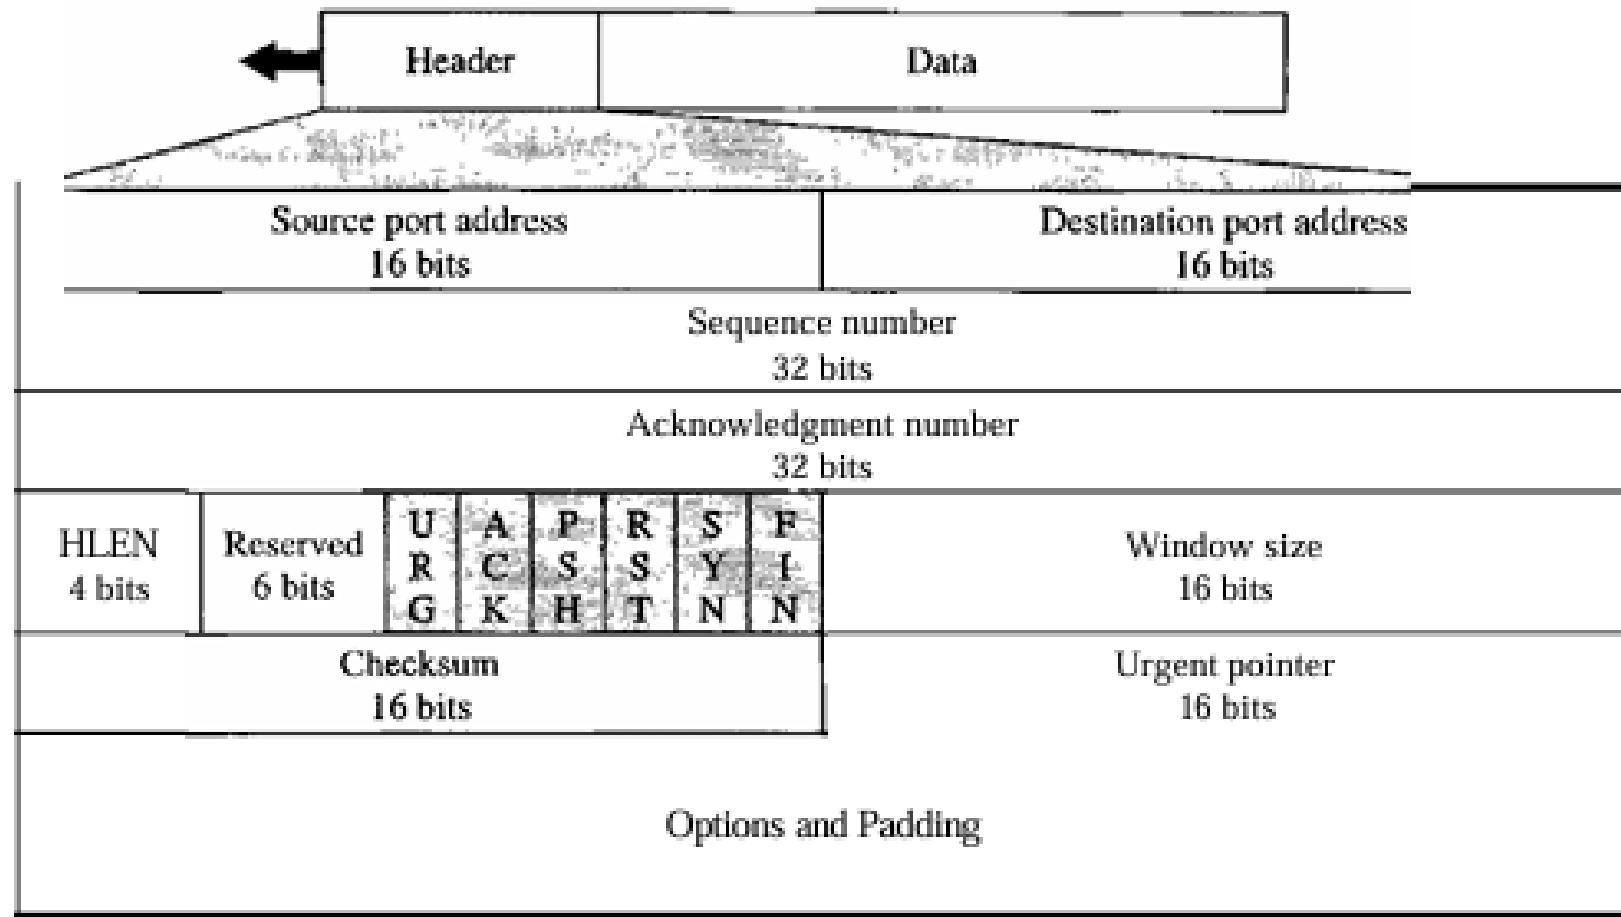


TCP Protocol: Flow Control

- **Flow Control:** TCP, Unlike UDP, provides flow control. The receiver of the data controls the amount of data that are to be sent by the sender. This is done to prevent the receiver from being over whelmed with data. The numbering system allows TCP to use a byte-oriented flow control.
- **Error Control:** To provide reliable service, TCP implements an error control mechanism. Although error control considers a segment as the unit of data for error detection (loss or corrupted segments), error control is byte-oriented.
- **Congestion Control:** TCP, unlike UDP, takes into account congestion in the network. The amount of data sent by a sender is not only controlled by the receiver (flow control), but is also detennined by the level of congestion in the network.

TCP Protocol Structure

Figure 23.16 *TCP segment format*



TCP Protocol Structure

- **Sequence number:** This 32-bit field defines the number assigned to the first byte of data contained in this segment. As we said before, TCP is a stream transport protocol. To ensure connectivity, each byte to be transmitted is numbered.
- **Acknowledgment number:** This 32-bit field defines the byte number that the receiver of the segment is expecting to receive from the other party. If the receiver of the segment has successfully received byte number x from the other party, it defines $x + 1$ as the acknowledgment number. Acknowledgment and data can be piggybacked together.
- **Header length:** This 4-bit field indicates the number of 4-byte words in the TCP header.

TCP Protocol Structure

- **Control:** This field defines 6 different control bits or flags as shown in Figure 23.17. One or more of these bits can be set at a time.

Figure 23.17 *Control field*

URG: Urgent pointer is valid

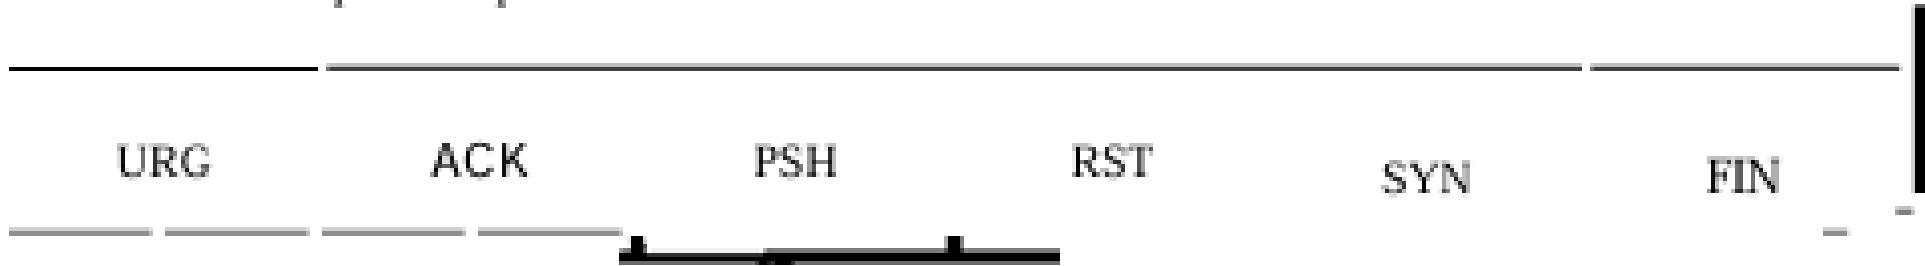
ACK: Acknowledgment is valid

PSH: Request for push

RST: Reset the connection

SYN: Synchronize sequence numbers

FIN: Terminate the connection



TCP Protocol Structure

- Control

Table 23.3 *Description of flags in the control field*

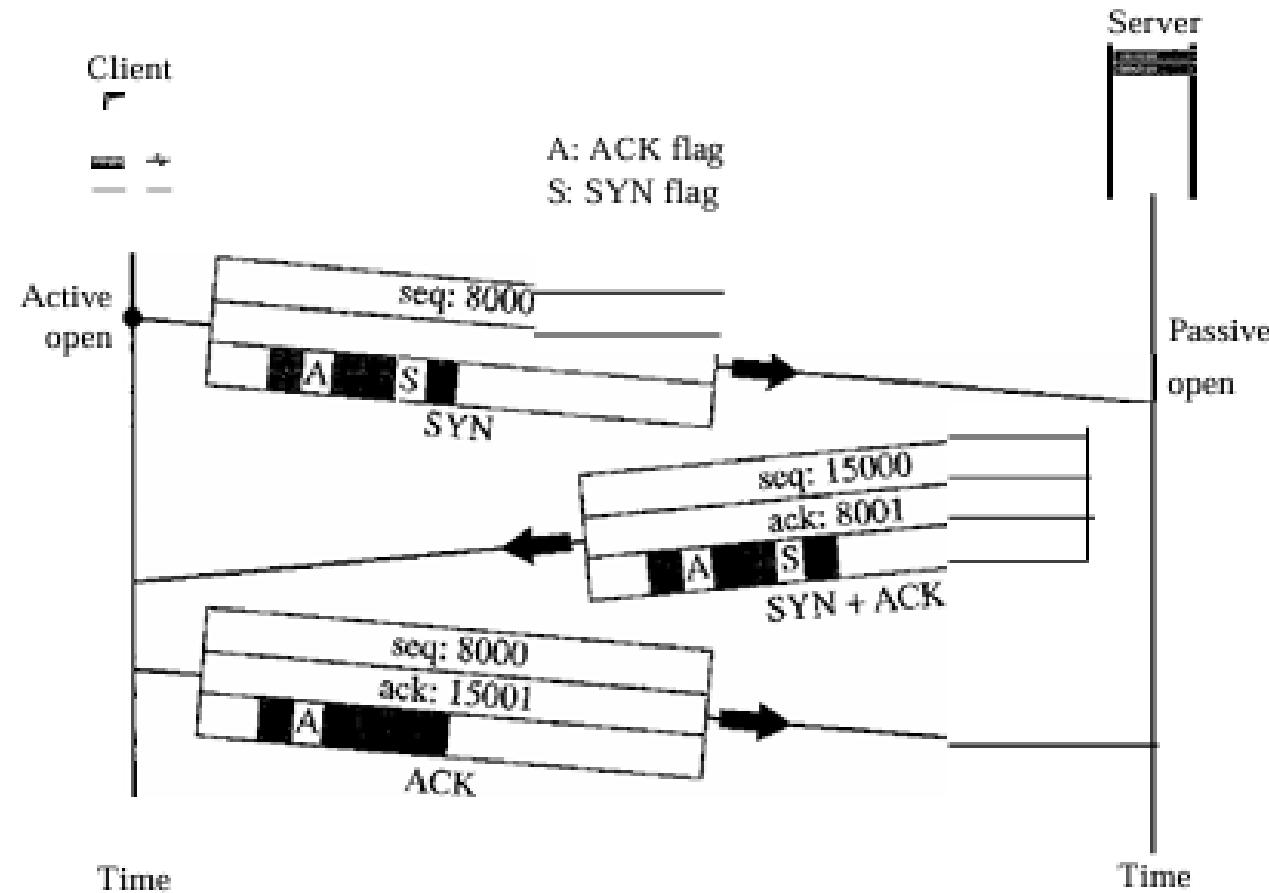
| <i>Flag</i> | <i>Description</i> |
|-------------|-------------------------------------------------|
| URG | The value of the urgent pointer field is valid. |
| ACK | The value of the acknowledgment field is valid. |
| PSH | Push the data. |
| RST | Reset the connection. |
| SYN | Synchronize sequence numbers during connection. |
| FIN | Terminate the connection. |

TCP Protocol Structure

- **Window size:** This field defines the size of the window, in bytes, that the other party must maintain. Note that the length of this field is 16 bits, which means that the maximum size of the window is 65,535 bytes. This value is normally referred to as the receiving window (rwnd) and is determined by the receiver. The sender must obey the dictation of the receiver in this case.
- **Urgent pointer:** This 16-bit field, which is valid only if the urgent flag is set, is used when the segment contains urgent data. It defines the number that must be added to the sequence number to obtain the number of the last urgent byte in the data section of the segment. Urgent data can be interrupt, abort etc.

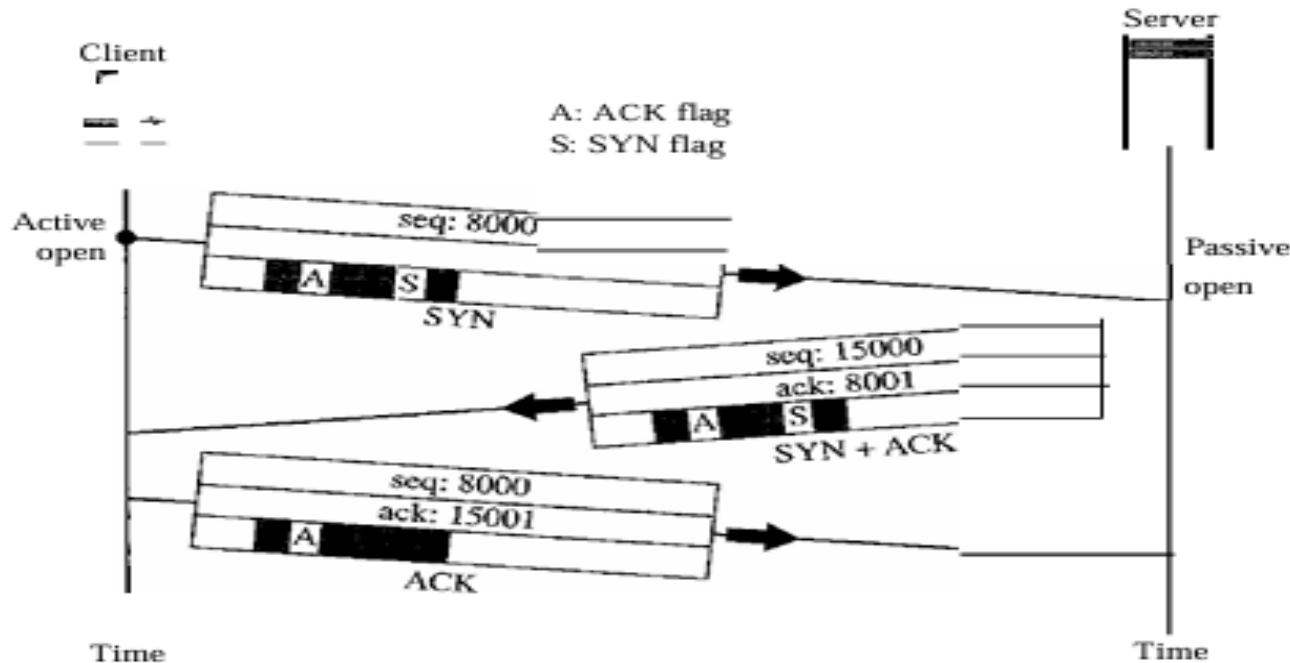
Connection Establishment in TCP Protocol

Figure 23.18 Connection establishment using three-way handshaking



Connection Establishment in TCP Protocol

Figure 23.18 Connection establishment using three-way handshaking



A SYN segment cannot carry data, but it consumes one sequence number.

A SYN + ACK segment cannot carry data,
but does consume one sequence number.

An ACK segment, if carrying no data, consumes no sequence number.

Connection Establishment in TCP Protocol

- **Solve Example 23.4 and 23.5**